

[Skip to main content](#)

- [Home](#)
- [Topics](#)
- [Reference](#)
- [Glossary](#)
- [Help](#)
- [Notebook](#)

Virtual Workshop

Welcome guest

[Log in \(Globus\)](#)

[Log in \(other\)](#)

[Try the quiz before you start](#)

Parallel I/O

[Overview](#) [Goals](#) [Prerequisites](#)

[The Parallel I/O Stack](#) [Lustre](#) • [Lustre Components](#) • [Lustre Interactions](#) • [Lustre Striping](#) • [File System Commands](#) • [A Striping Test to Try](#) [Parallel I/O with MPI-IO](#) • [MPI-IO Advantages](#) • [File Pointers and Offsets](#) • [File Views](#) • [File View Examples](#) • [Collective I/O](#) • [Collective Choreography](#) • [Asynchronous Operations](#) • [Summary of Variants](#) • [Passing Along Hints](#) • [MPI-IO Summary](#) [Higher-Level Alternatives](#) [Exercise Quiz](#) [Short survey](#)

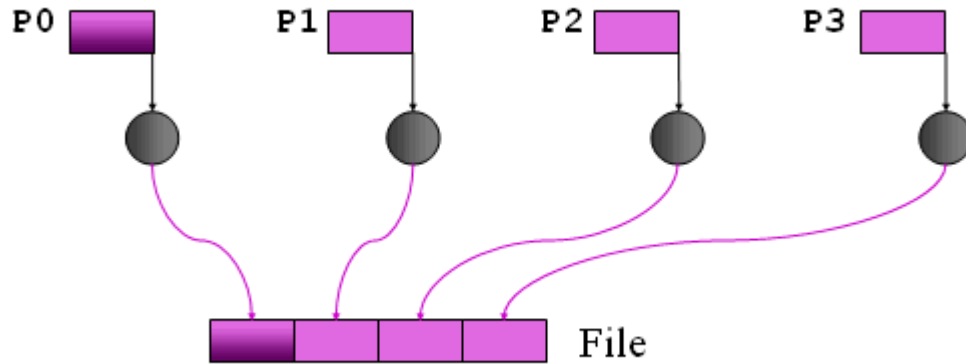
File View Examples

Example 1: We write contiguous data into a contiguous block defined by a file view. We give each process a different file view using `MPI_File_set_view` so that together, the processes lay out a series of blocks in the file, one block per process.

```
#define N 100
MPI_Datatype arraytype;
MPI_Offset disp;

disp = rank*sizeof(int)*N; etype = MPI_INT;
MPI_Type_contiguous(N, MPI_INT, &arraytype);
MPI_Type_commit(&arraytype);

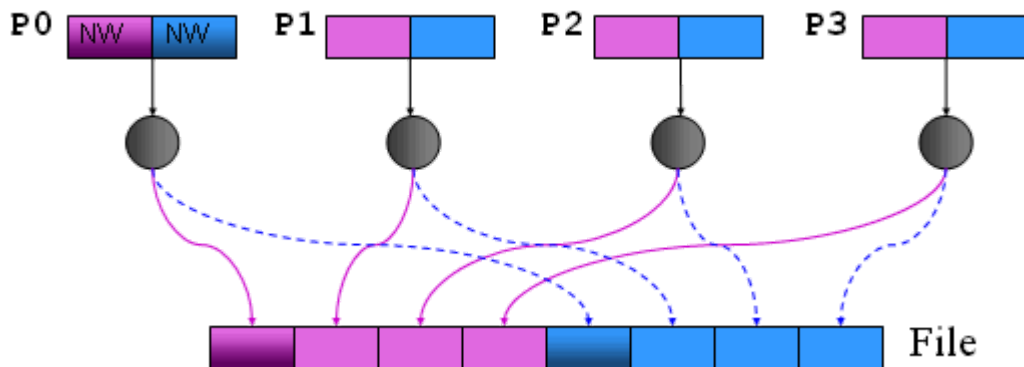
MPI_File_open(    MPI_COMM_WORLD, "/pfs/datafile",
                  MPI_MODE_CREATE | MPI_MODE_RDWR,
                  MPI_INFO_NULL, &fh);
MPI_File_set_view(fh, disp, etype, arraytype,
                  "native", MPI_INFO_NULL);
MPI_File_write(fh, buf, N, etype, MPI_STATUS_IGNORE);
```



Example 2: We write contiguous data into two *separate* blocks defined by a file view. Each block is a contiguous type in memory, but the pair of blocks is a vector type in the file view. We again use displacements to lay out a series of blocks in the file, so that two blocks are written per process, in a repeating fashion.

```
int buf[NW*2];
MPI_File_open(MPI_COMM_WORLD, "/data2",
              MPI_MODE_RDWR, MPI_INFO_NULL, &fh);
/* want to see 2 blocks of NW ints, NW*npes apart */
MPI_Type_vector(2, NW, NW*npes, MPI_INT, &fileblk);
MPI_Type_commit(&fileblk);
disp = (MPI_Offset)rank*NW*sizeof(int);
MPI_File_set_view(fh, disp, MPI_INT, fileblk,
                 "native", MPI_INFO_NULL);

/* processor writes 2 'ablk', each with NW ints */
MPI_Type_contiguous(NW, MPI_INT, &ablk);
MPI_Type_commit(&ablk);
MPI_File_write(fh, (void *)buf, 2, ablk, &status);
```



- The data type in memory is just a contiguous set of NW ints.
- The file view is of type fileblk, displaced by rank *NW*sizeof(int).
- Two units of the data type are written into the file whose view has been set.

[<= previous](#)

[switch](#)

[next =>](#)

Add my notes

Mark (M) my place in this module

